Utilitarianism

ACT UTILITARIANISM

Jeremy Bentham defended the 'principle of utility' or 'greatest happiness principle', 'that principle which approves or disapproves of every action whatsoever, according to the tendency which it appears to have to augment or diminish the happiness of the party whose interest is in question'. Or again, 'that principle which states the greatest happiness of all those whose interest is in question, as being the right and proper… end of human action'. (*Introduction to the Principles of Morals and Legislation*, Ch. 1 §2)

If we simplify this a little, we can say that act utilitarianism claims that an action is right if it leads to the greatest happiness of all those it affects, i.e. if it maximizes happiness. Otherwise, the action is wrong. The greatest happiness should be the goal of our actions, what we hope to bring about. Our actions are judged not 'in themselves', e.g. by what type of action they are (a lie, helping someone, etc.), but in terms of what consequences they have.

'Greatest happiness' is comparative (great, greater, greatest). If an action leads to the greatest happiness of those it affects, no other action taken at that time could have led to greater happiness. So an action is right only if, out of all the actions you could have done, this action leads to more happiness than any other. Just causing some happiness, or more happiness than unhappiness, isn't enough for an act to be morally right.

Act utilitarianism seems to have the advantage of providing a clear and simple way of making decisions – consider the consequences of the different actions you could perform and choose that action that brings about the greatest happiness. It makes complicated decisions easy – the only thing that matters is happiness (and surely everyone wants happiness).

Objections

But is the guidance offered is helpful? How can we know or work out the consequences of an action, to discover whether it maximizes happiness or not? Bentham does not say that an action is right if it actually maximizes happiness. He says it is right according to 'the tendency which it appears to have' to maximize happiness. We don't need to be able to work out the consequences precisely. An action is right if we can reasonably expect that it will maximize happiness.

This still means we must be able to work things out roughly. John Stuart Mill, who defended a different version of utilitarianism, thought this was still too demanding. Happiness is 'much too complex and indefinite' a standard to apply directly to actions. But we don't need to try, he claimed, because over time, people have automatically, through trial and error, worked out which actions tend to produce happiness. This is what our inherited moral rules actually are: 'tell the truth', 'don't steal', and 'keep your promises' are embodiments of the wisdom of humanity that lying, theft, and false promising tend to lead to unhappiness. Mill called these moral rules 'secondary principles'. It is only in cases of conflict between secondary principles (e.g. if by telling

the truth you break your promise) that we need to apply the greatest happiness principle directly to an action.

A second criticism of act utilitarianism is that no type of action is ruled out as immoral. If torturing a child produces the greatest happiness, then it is right to torture a child. Suppose a group of child abusers only find and torture abandoned children. Only the child suffers pain (no one else knows about their activities). But they derive a great deal of happiness. So more happiness is produced by torturing the child than not, so it is morally right. This is clearly the wrong answer.

Act utilitarians can reply that it is very probable that someone will find out, and then many people will be unhappy. Because we should do what is likely to produce the greatest happiness, we shouldn't torture children. However, the theory still implies that if it was very unlikely anyone would find, then it would be right to torture children. But other people finding out isn't what makes torturing children wrong.

Integrity and demandingness
Act utilitarianism says that my happiness doesn't count any more than anyone else's when I'm considering what to do. Obviously, I am affected more often and more deeply by my actions than other people. But that's it. My actions is just a means of generating the greatest overall happiness.

So what if my conscience and moral integrity tell me not to do something that act utilitarianism says it is right to do? Act utilitarianism says I am supposed to calculate how unhappy it would be make me to perform the action, and see if that changes what I should do! If it doesn't, I should do the action anyway. This doesn't 'respect' the importance of people's integrity.

Second, act utilitarianism is too demanding. For example, every time I buy a CD, I could have given the money to charity; and surely that would create more happiness, since other people need food more than I need music. But because people will always need food, it will never be right for me to buy myself music. It will never be right to do something just for myself if I have more than the bare minimum.

RULE UTILITARIANISM
Rule utilitarianism claims that an action is right if, and only if, it complies with those rules which, if everybody followed them, would lead to the greatest happiness (compared to any other set of rules). (By contrast, Mill's secondary principles are rules of thumb, i.e. not strict rules that we must follow, but helpful guidance in our thoughts about what to do.)

Rule utilitarianism has some advantages over act utilitarianism. First, the rule forbidding torture of children will clearly cause more happiness if everyone followed it than the rule allowing torture of children. So it is wrong to torture children. Second, we don't have to work out the consequences of each act in turn to see if it is right. We need to work out which rules create the greatest happiness, but we only need to do this once, and we can do it together.

Third, the rule that we should usually allow people to act on integrity is a rule that will promote more happiness than any other. Finally, morality is not so demanding: I am only

required to act in a way that, if everyone acted like that, would promote the greatest happiness. In the case of charity, I only need to give as much to charity as would be a 'fair share' of the amount needed to really help other people. In all these ways, it looks like rule utilitarianism provides better guidance for making moral decisions.

However, act utilitarians object that rule utilitarianism amounts to 'rule-fetishism'. The point of the rules is to bring about the greatest happiness. If we only give as much to charity as we would need to if everyone gave to charity, then many people will not be helped, because not everyone will give what they should to charity. Surely, knowing this, I ought to give much more to charity; spending the money on myself would not be right. Or again, if I know that, e.g. lying in a particular situation will produce more happiness than telling the truth, it seems pointless to tell the truth, causing unhappiness. The whole point of the rule was to bring about happiness, so there should be an exception to the rule in this case. But then whenever a particular action causes more happiness by breaking the rule than by following it, we should do that action. And then we are back with act utilitarianism, weighing up the consequences of each action in turn.

Rule utilitarians respond, first, that only rule utilitarianism can provide real guidance for making moral decisions; and second, that deciding according act utilitarianism will break down our trust that people behave morally. So following a rule even when in the particular case it will cause less happiness than breaking the rule, is still justified, because if people kept breaking the rules, that would cause less happiness in the long run.

VARIATIONS ON 'HAPPINESS'
Hedonism
If we are to bring about happiness, we should consider what happiness is. Bentham and Mill were 'hedonists', claiming that happiness is pleasure and the absence of pain, and that this is all that matters. Bentham argued that we can measure pleasures and pains and add them up on a single scale by a process he called the 'felicific calculus'. If a pleasure is more intense, will last longer, is more certain to occur, will happen sooner rather than later, or will produce in turn many other pleasures and few pains, it counts for more. In thinking what to do, you also need to take into account how many people will be affected (the more you affect positively, and the fewer you affect negatively, the better). The total amount of happiness produced is the sum total of everyone's pleasures produced minus the sum total of everyone's pains.

Mill rejected Bentham's view that pleasures and pains are all equally important. Some types of pleasure are 'higher' than others, more valuable, more important to human happiness. But which ones? Mill thought there was an objective test: if almost everyone who knows what they are talking about compares two pleasures and agrees that the first is 'more desirable and valuable' than the second, then the first is a higher pleasure.

But how can we tell if a pleasure is more valuable (quality) than another, rather than just more pleasurable (quantity)? To tell if a pleasure is more valuable, people have to prefer it even if having that pleasure brings more pain with it. For example, the pleasure of being in love carries the pain of longing and the possible pain of breaking up. But people still prefer being in love to, say, eating delicious food. Happiness is not contentment.

Mill argues that, as long as our physical needs are met, people will prefer the pleasures of thought, feeling, and imagination to pleasures of the body and the senses, even though

our 'higher' capacities also mean we can experience terrible pain, boredom, dissatisfaction.

Pleasure and preferences
Preference utilitarianism claims we should maximize not pleasure, but the satisfaction of people's preferences. Two reasons favour this theory. First, it is easier to know whether someone's preference has been satisfied than how much pleasure someone experiences. So it provides better guidance. Second, it can be right to satisfy someone's preferences even when they don't know this has happened, and so don't derive any pleasure from it. For example, I can want you to look after my ant farm when I die. You should still look after my ants, rather than kill them, even though I cannot gain any pleasure for this.

Objections
But is happiness the right standard for morality at all? Utilitarianism weighs the unhappiness of one person against the happiness of another, whether this is in deciding which action to do or which rule to adopt. Philosophers object that it isn't concerned with people as individuals, but as 'receptacles' for happiness, which fails to show them proper respect. Furthermore, the distribution of happiness – who gets happy by how much – is irrelevant, which fails to respect justice.

Second, Kant argued that happiness (or satisfying people's preferences) is not always morally good. For example, the happiness child abusers get from hurting children is morally bad. The fact that the abusers are made happy by what they do doesn't make their action better at all, but worse. So there must be some other standard than happiness for what is morally good.

Third, is happiness the only thing that matters? Many philosophers argue that there are other values (freedom, justice…). Even if pursuing these values gives us happiness, it is not the happiness the value brings which is important, but the value itself.